



PULP
Parallel Ultra Low Power

PULP RISC-V Cores: Ariane, RI5CY and friends

HPCA 2018 - Vienna

25.02.2018



Florian Zaruba

Frank K. Gürkaynak

Andreas Kurth

Francesco Conti



Multitherman

 **PRECOMP**
Open Transprecision Computing

 **ExaNode**

¹*Department of Electrical, Electronic
and Information Engineering*

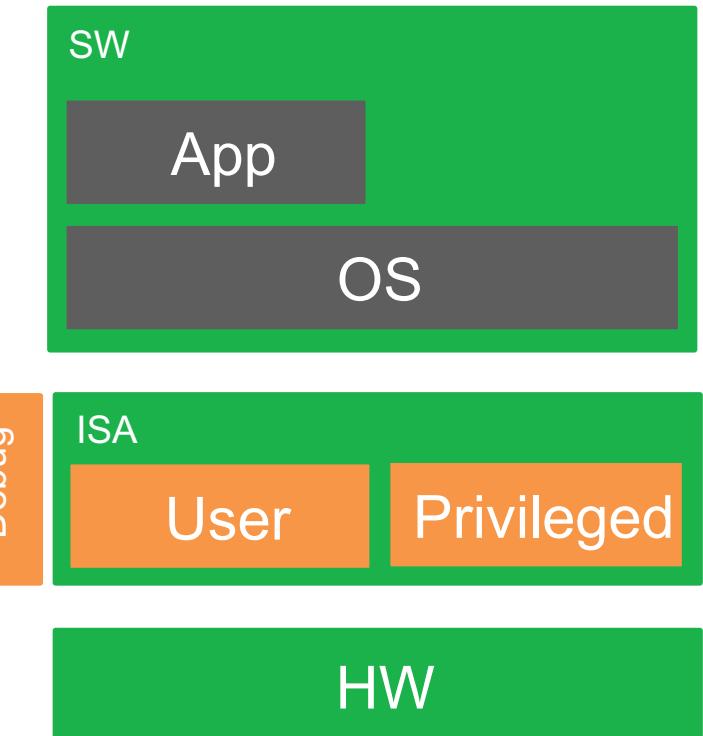
ETH zürich

²*Integrated Systems Laboratory*

<http://pulp-platform.org>

RISC-V Instruction Set Architecture

- Started by UC-Berkeley in 2010
- Contract between SW and HW
 - Partitioned into user and privileged spec
 - External Debug
- Open Standard governed by RISC-V foundation
 - ETHZ is a founding member of the foundation
 - Necessary for the continuity
- Defines 32, 64 and 128 bit ISA
 - No implementation, just the ISA
 - Different RISC-V implementations (both open and close source) are available
- At IIS we specialize in
efficient implementations of RISC-V cores



RISC-V ISA

- Generally kept very simple and extendable
- From deeply-embedded devices to HPC
- ISA support is given by RV + word-width + extensions supported (RV32I)
- User specification:
 - Separated into extensions, only I is mandatory
- Privileged Specification (WIP):
 - Governs OS functionality: Exceptions, Interrupts
 - Virtual Addressing
 - Privilege Levels



Spec separated into “extensions”

- | | |
|---|--|
| I | Integer instructions (frozen) |
| E | Reduced number of registers |
| M | Multiplication and Division (frozen) |
| A | Atomic instructions (frozen) |
| F | Single-Precision Floating-Point (frozen) |
| D | Double-Precision Floating-Point (frozen) |
| C | Compressed Instructions (frozen) |
| X | Non Standard Extensions |

RISC-V Planned Extensions

- No or partial work done yet on those extensions
- Possible to contribute as a foundation member in **task-groups**
- Dedicated task-groups
 - Formal specification
 - Memory Model
 - Marketing
 - External Debug Specification
- For Bit-manipulation we provide our own solution → part of the task group



Q Quad-Precision Floating-Point

L Decimal Floating-Point (IEEE 754-2008)

B Bit-Manipulation

T Transactional Memory

P Packed-SIMD

J Dynamically Translated Languages

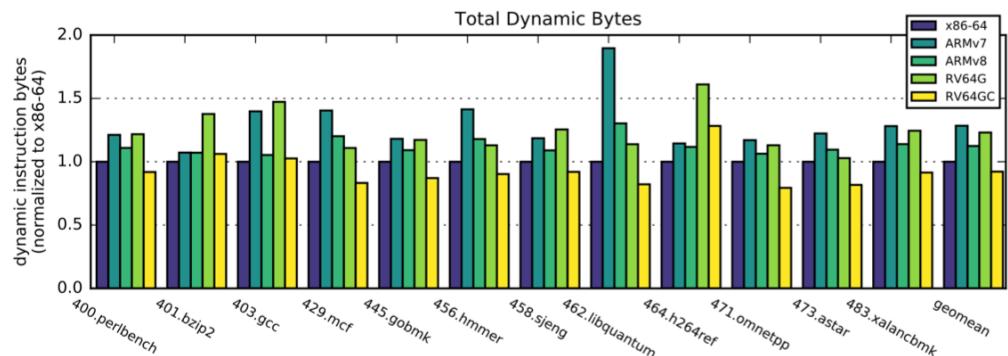
V Vector Operations

N User-Level Interrupts

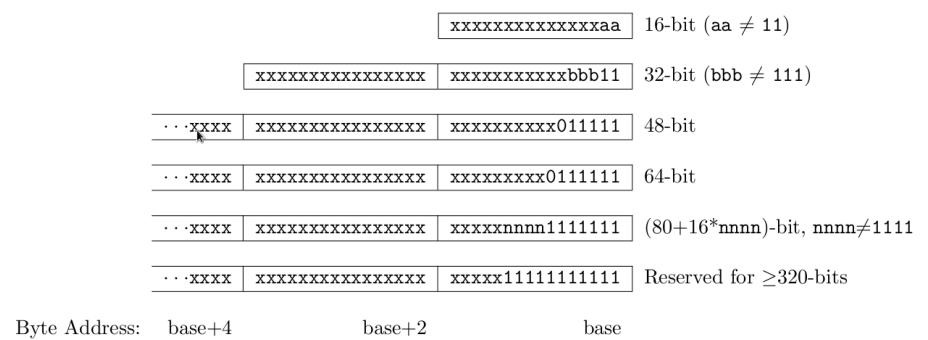
Compressed Instructions “C”



- 16-bit instructions covering most used ops, code size reduction by 34 %
- Length-encoded in LSBs, variable length instructions
- Compressed instructions increase fetch-bandwidth
- Allow for macro-op fusion of common patterns

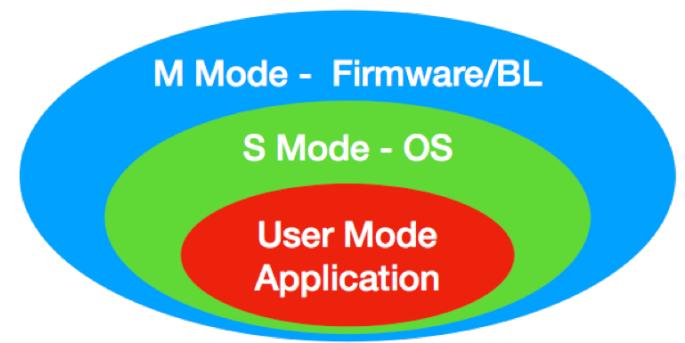


x86-64: 3.71 bytes / instruction **RV64IC:** 3.00 bytes / instruction



Privileged Specification 1.11 (WIP)

- Defines Control and Status Registers (CSR)
- Defines instructions to RMW CSR
- 3 Rings of operation
 - Machine-, Supervisor- and User-Mode
 - Hypervisor WIP
- Exceptions + IRQ support
 - IRQ/Exception stack
 - Wait for Interrupt (WFI) instruction
 - Specification for platform level interrupt controller
 - Instructions to enter/return from exceptions
- Virtual Memory
 - Page-based, 32-bit PA, 39-bit PA and/or 48-bit PA
- Platform Configuration String (DTS)



RISC-V Ecosystem¹



- Binutils – upstream
- GCC – upstream
- LLVM – in development
- Simulator:
 - "Spike" - reference
 - QEMU, Gem5
- OpenOCD
- OS: Linux, sel4, freeRTOS, zephyr
- Runtimes: Jikes, Ocaml, Go
- SW maintained by different parties
 - Binutils and GCC by Sifive a Berkeley start-up

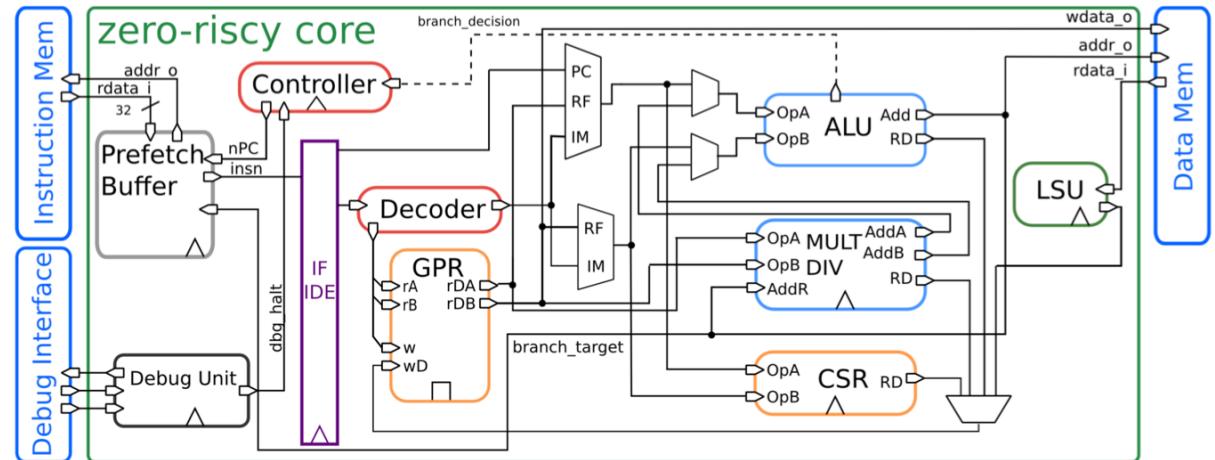
¹ <https://github.com/riscv/riscv-wiki/wiki/RISC-V-Software-Status>

RISC-V cores under development at IIS

32 bit			64 bit
Low Cost Core	Core with DSP enhancements	Floating-point capable Core	Linux capable Core
<ul style="list-style-type: none">■ Zero-riscy<ul style="list-style-type: none">■ RV32-ICM■ Micro-riscy<ul style="list-style-type: none">■ RV32-CE <p>ARM Cortex-M0+</p>	<ul style="list-style-type: none">■ RI5CY<ul style="list-style-type: none">■ RV32-ICMX■ SIMD■ HW loops■ Bit Man■ Fixed point <p>ARM Cortex-M4</p>	<ul style="list-style-type: none">■ RI5CY + FPU<ul style="list-style-type: none">■ RV32-ICMFX <p>ARM Cortex-M4F</p>	<ul style="list-style-type: none">■ Ariane<ul style="list-style-type: none">■ RV64-IC(MA)■ Full privileged specification <p>ARM Cortex-A5</p>

Zero/Micro-riscy, small area core for control applications

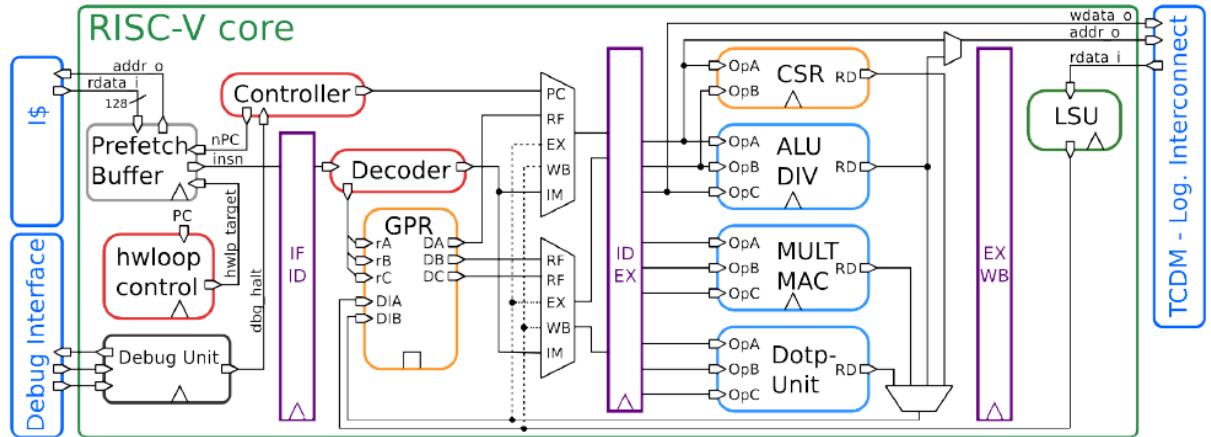
- 2-stage pipeline
- Optimized for area
 - Area:
 - 19 kGE (Zero-riscy)
 - 12 kGE (Micro-riscy)
 - Critical path:
 - ~ 30 logic levels
- Plans for tape-out
 - Early version taped out in UMC65
 - Controller core of Mr. Wolf (PULP systems in TSMC40 LP)



- Two Configurations:
 - **Zero-riscy**: RV32-ICM (2,44 Coremark/MHz)
 - 32 registers, hardware multiplier
 - **Micro-riscy** : RV32-CE (0,91 Coremark/MHz)
 - 16 registers (E), software emulated multiplier

RI5CY: Our 32-bit workhorse

- **4-stage pipeline**
 - 41 kGE
 - 30 logic levels of critical path
 - Coremark/MHz 3.19
- **Includes various extensions**
 - SIMD
 - Fixed point
 - Bit manipulations
 - HW loops
- **Versions taped out in:**
 - GF28
 - UMC65
 - TSMC40LP (with FPU, in preparation)



- **Different Options:**
 - **FPU:** IEEE 754 single precision
 - Including hardware support for FDIV, FSQRT, FMAC, FMUL
 - **Privilege support:**
 - Supports privilege mode **M** and **U**

A new perspective: Application class processor

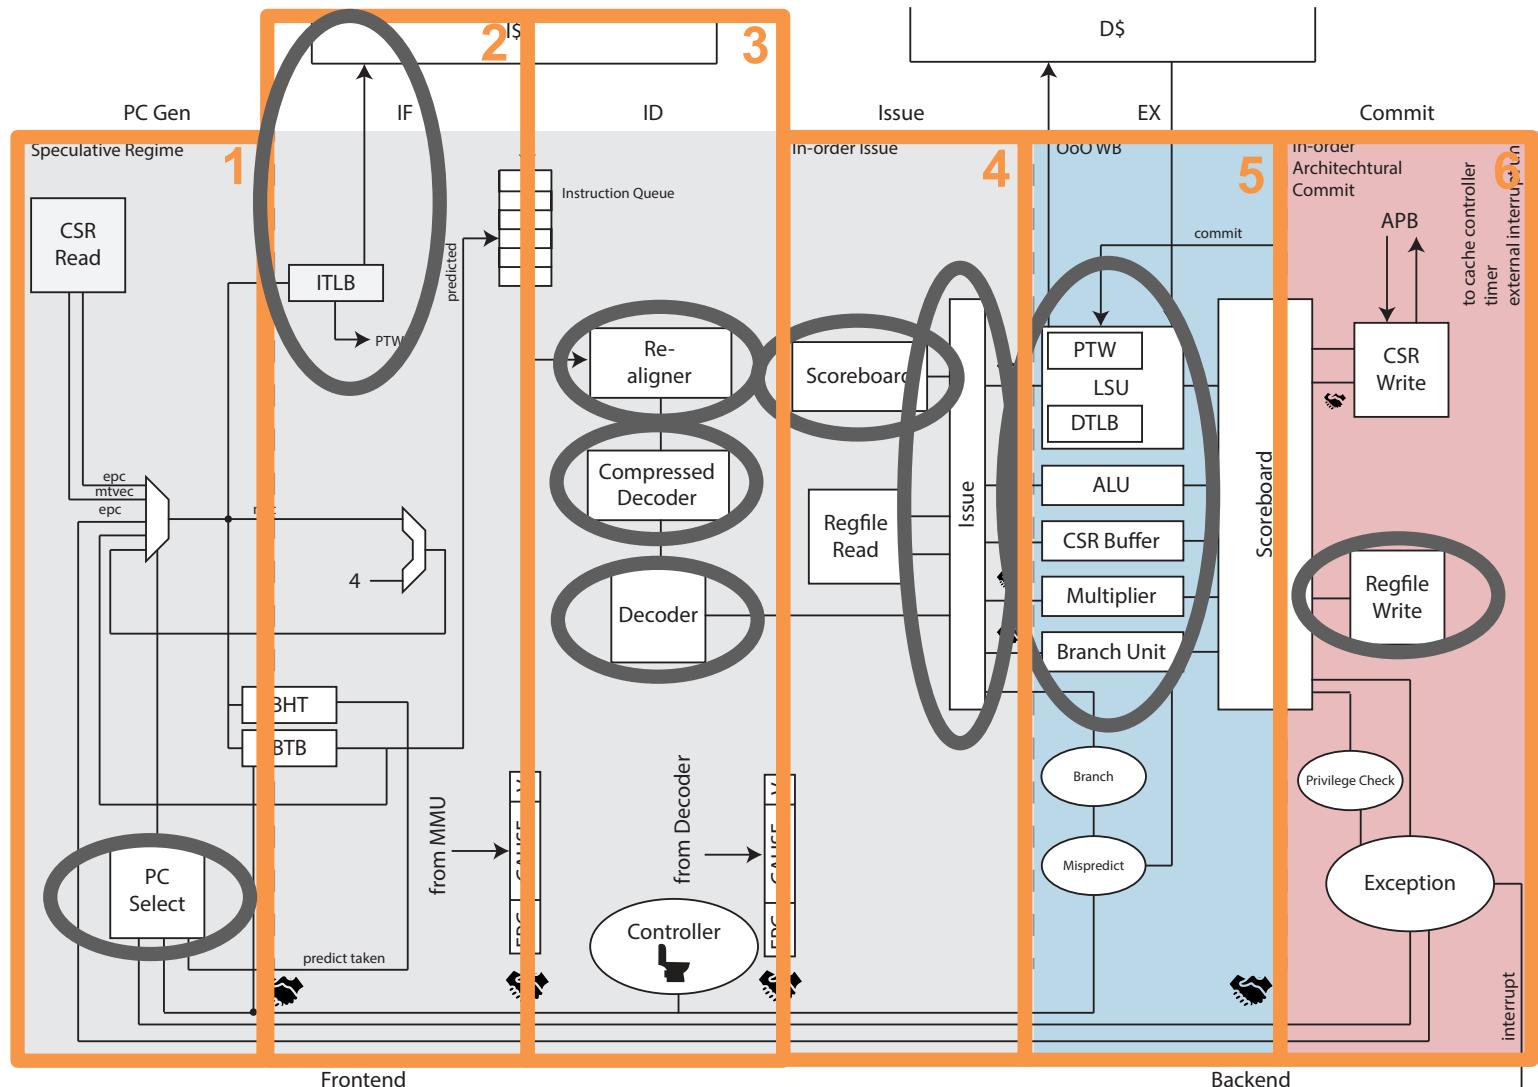
- Currently: Bare metal/simple run-time
 - Virtual Memory
 - Multi-program environment
 - Efficient sharing and protection
 - Operating System
 - Highly sequential code
 - Increase frequency to gain performance
 - Large software infrastructure
 - Drivers for hardware (PCIe, ethernet)
 - Application SW (e.g.: Tensorflow, ...)
 - Larger address space (64-bit)
 - **Requires more hardware support**
 - MMU (TLBs, PTW)
 - Privilege Levels
 - More Exceptions (page fault, illegal access)
- **Ariane** an application class processor



ARIANE: Linux Capable 64-bit core

- Application class processor
- Linux Capable
 - M, S and U privilege modes
 - TLB
 - Tightly integrated D\$ and I\$
 - Hardware PTW
- Optimized for performance
 - Frequency: > 1.5 GHz (22 FDX)
 - Area: 185 kGE
 - Critical path: ~ 25 logic levels
- 6-stage pipeline
 - In-order issue
 - Out-of-order write-back
 - In-order commit
- Branch-prediction
- Scoreboarding
- Taped-out in GF22FDX
- Designed for extensibility

- 1. PC Gen**
 - Select PC
- 2. Instr. Fetch**
 - TLB
 - Query I\$
- 3. Instr. Decode**
 - Re-align
 - De-compress
 - Decode
- 4. Issue**
 - Select FU
 - Issue
- 5. Execute**
- 6. Commit**
 1. Write state



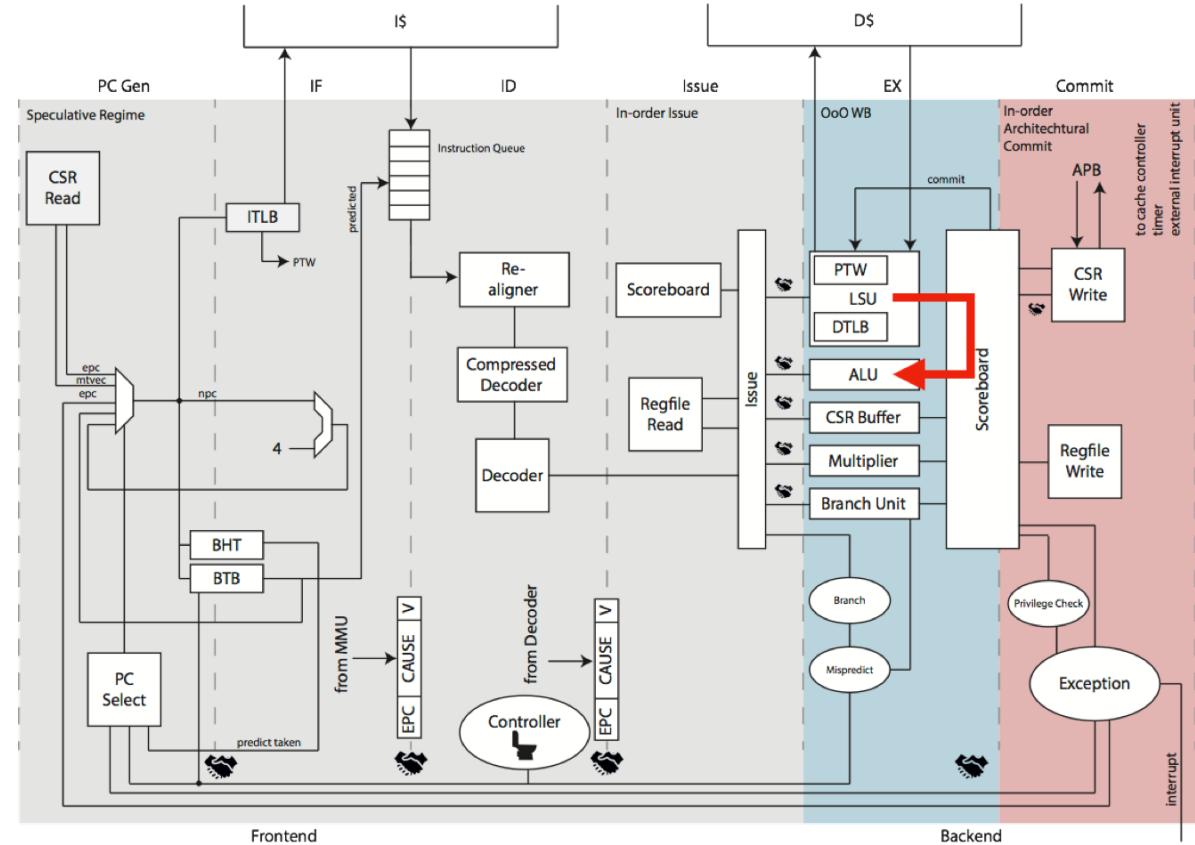
Frequency-IPC trade-off

- Frequency:
 - Increase frequency through pipelining
 - Modern Intel CPUs have around 10 - 20 pipeline stages
- Adds significant complexity on the cache interfaces
- Increased bubbles due to:
 - Data Hazards → **Forwarding**
 - Structural Hazards → **Scoreboard**
 - Control Hazards → **Branch Prediction**

Data Hazards - Forwarding

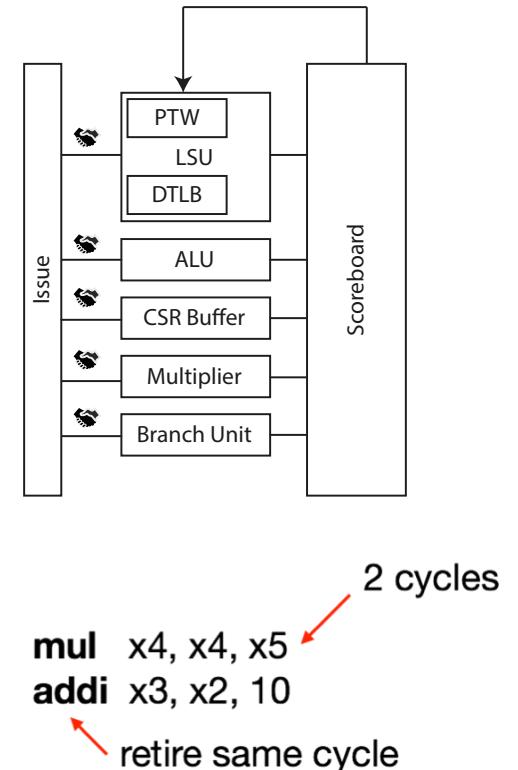
Data Hazards → Forwarding

Id x2, 0(x7)
addi x3, x2, 10



Scoreboarding

- Hide latency of multi-cycle instructions
- Clean and modular interface to functional units
→ scalability (FPU)
- Add issue port: Dual-Issue implementation
- Split execution into four steps:
 - **Issue**: Relatively complex issue logic (extra pipeline-stage)
 - **Read Operands**: From register file or forwarded
 - **Execute**
 - **Write Back**: Mitigate structural hazards on write-back path
- Implemented as a circular buffer



Scoreboard

Decode:

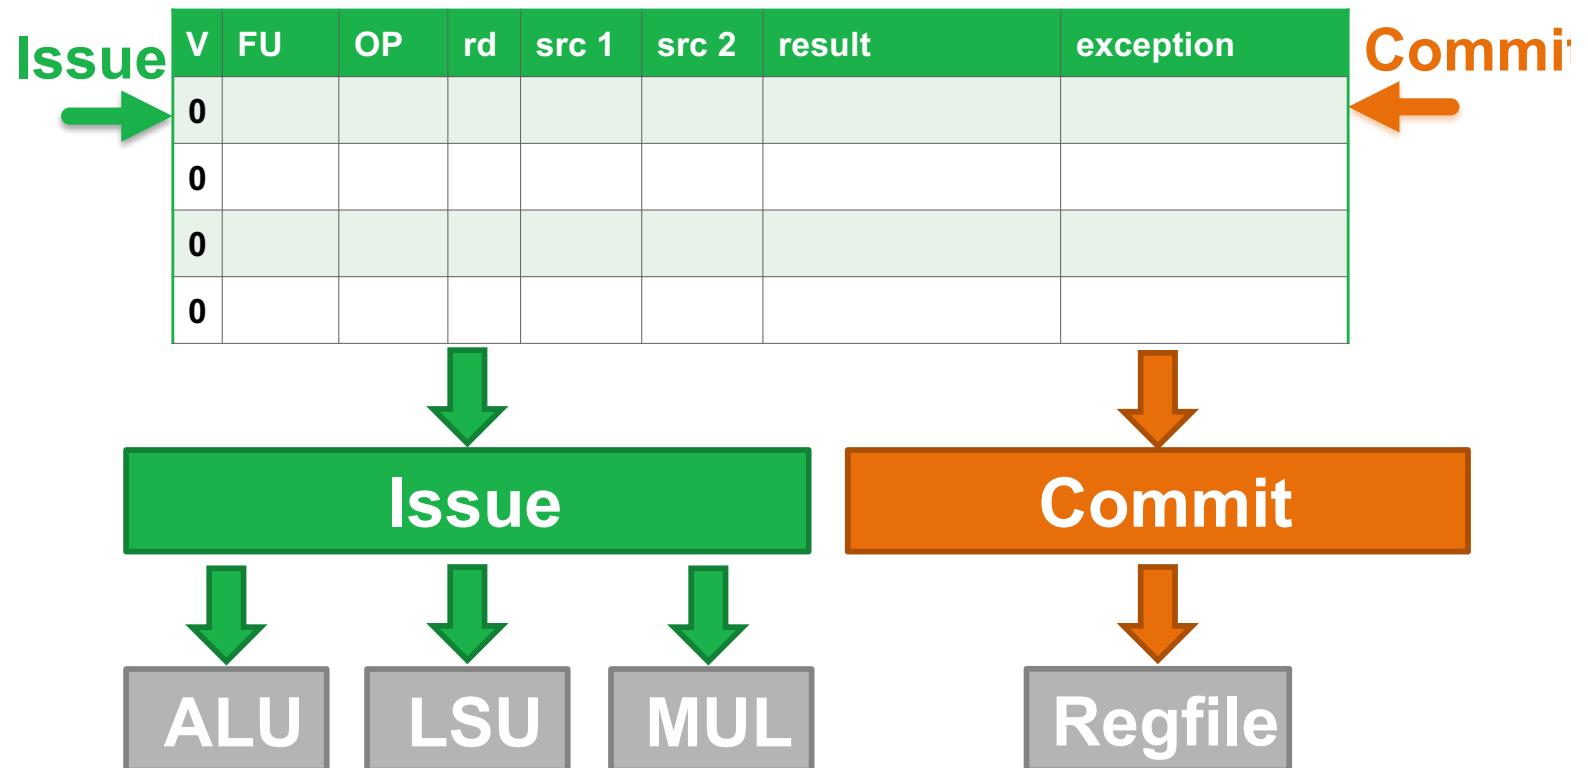
ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

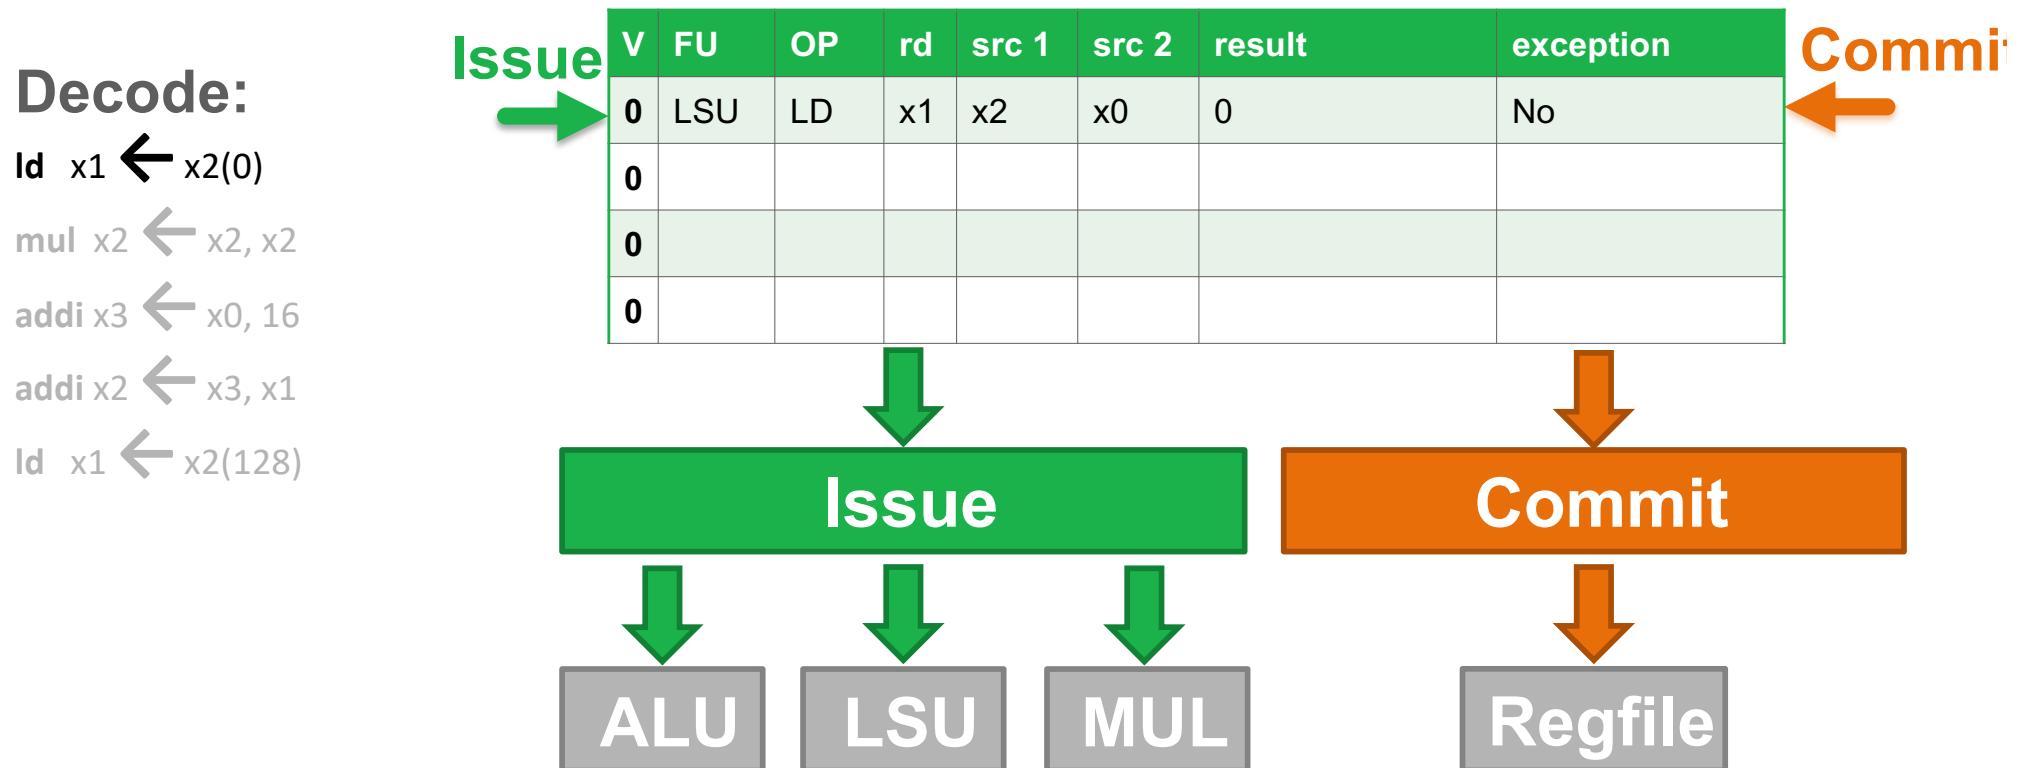
addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)



Scoreboard – 3 cycle instruction (LD)



Scoreboard – 2 cycle instruction (MUL)

Decode:

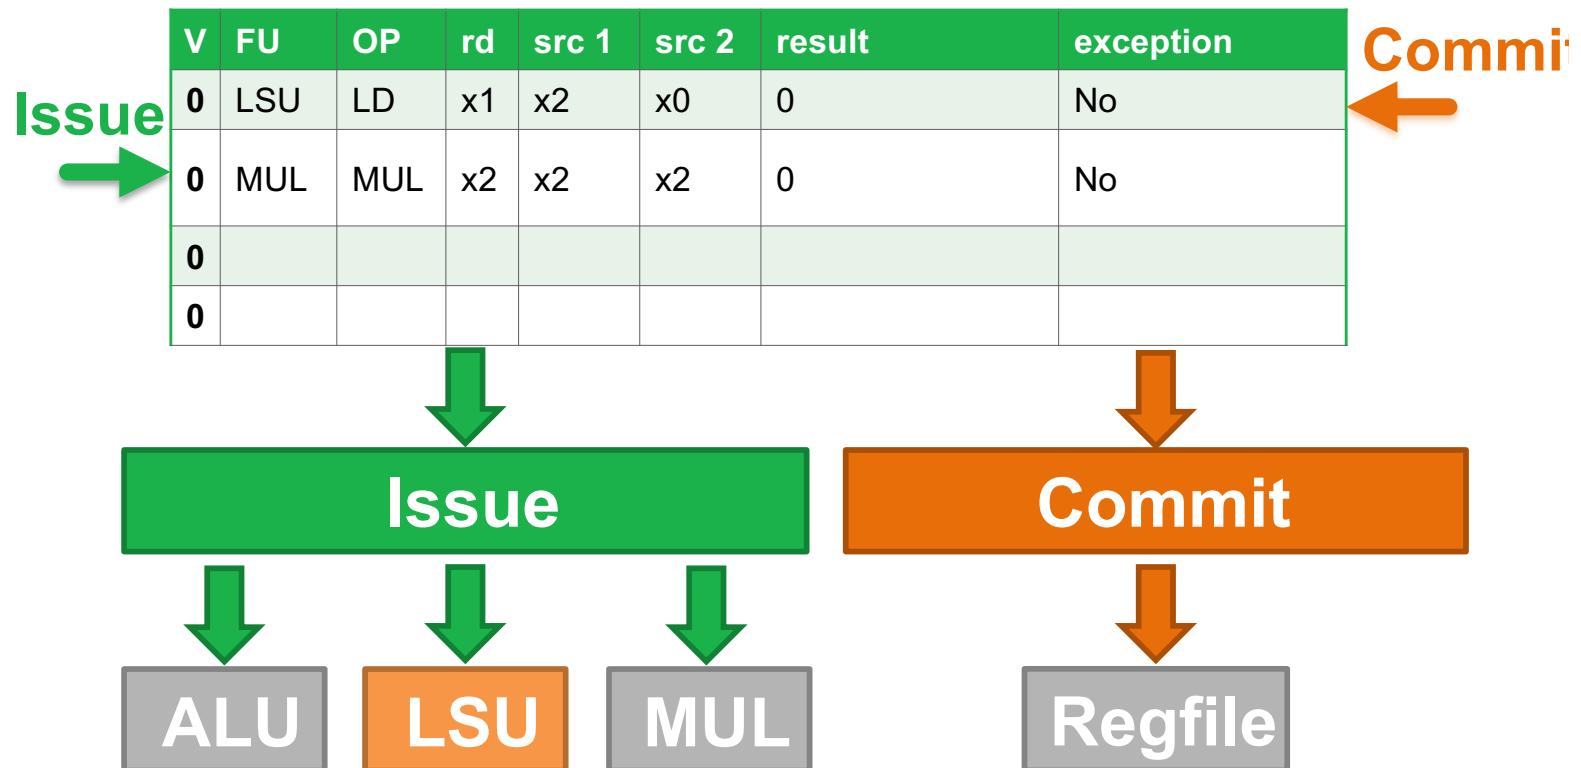
ld x1 ← x2(0)

mul x2 ← x2, x2

addi x3 ← x0, 16

addi x2 ← x3, x1

ld x1 ← x2(128)



Scoreboard - single cycle instruction (ALU)

Decode:

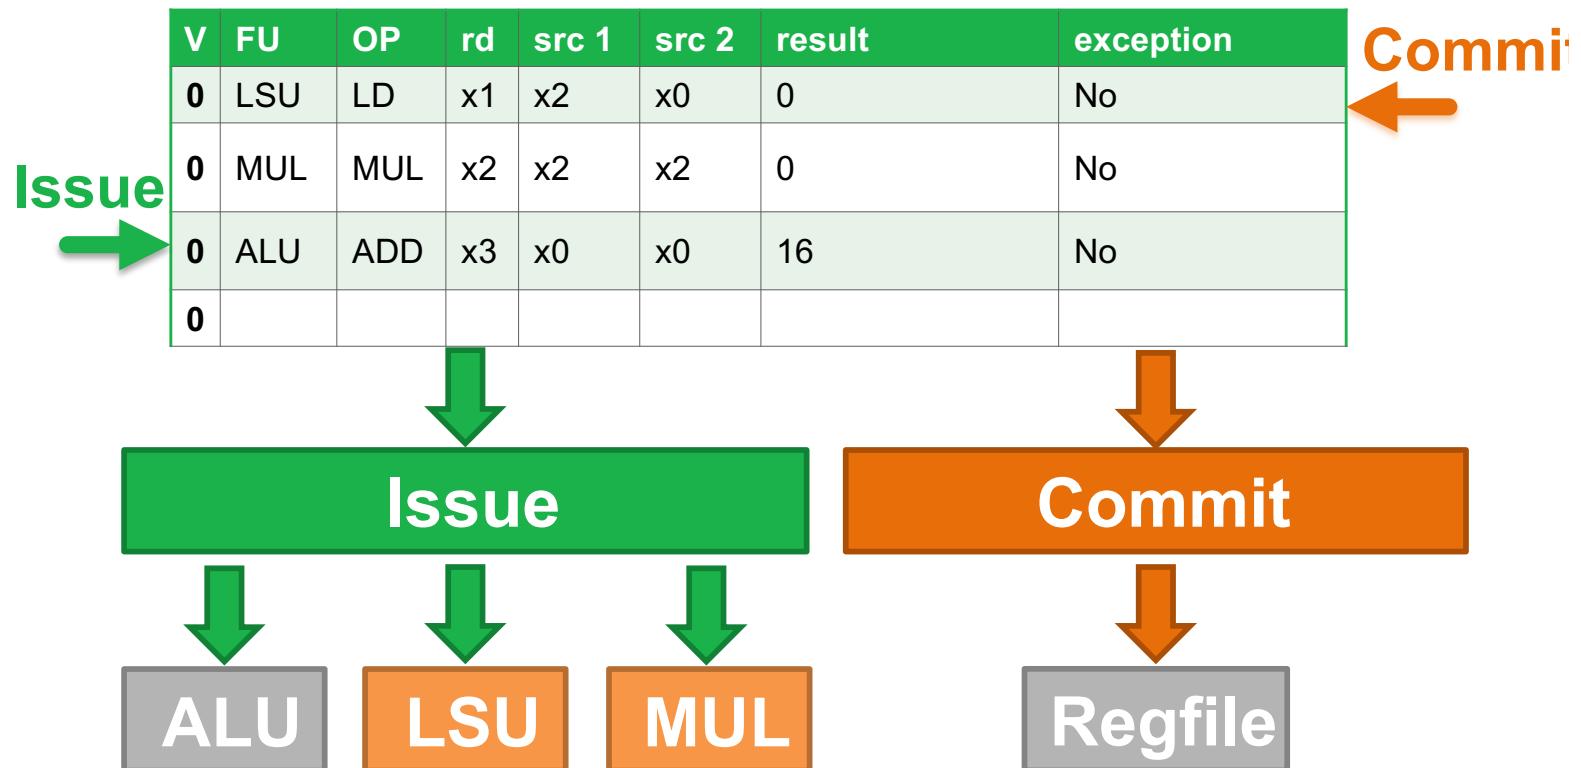
ld x1 ← x2(0)

mul x2 ← x2, x2

addi x3 ← x0, 16

addi x2 ← x3, x1

ld x1 ← x2(128)



Scoreboard – Multiple Write Back

Decode:

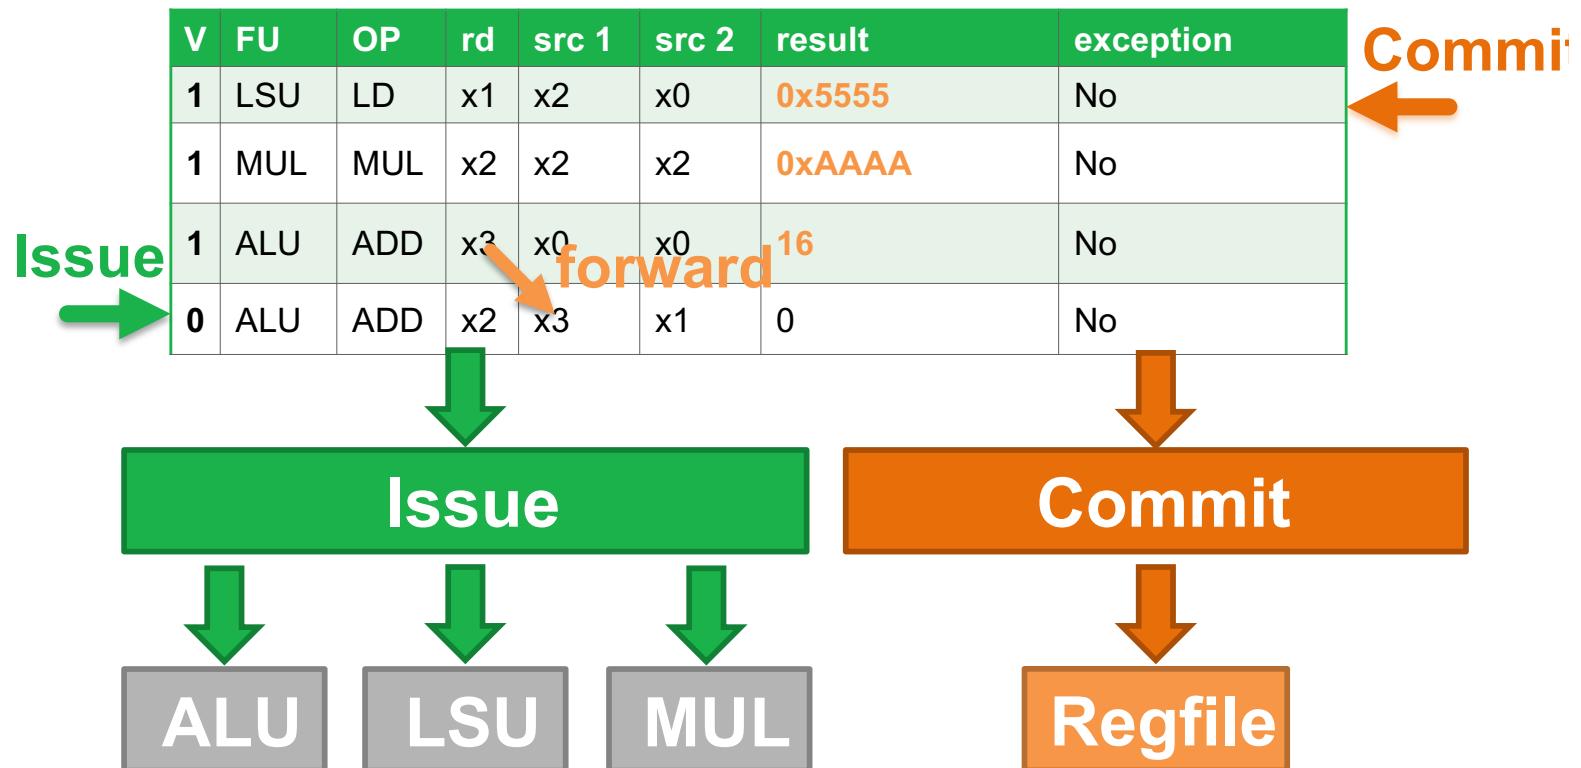
ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)



Scoreboard - Commit

Decode:

ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)

Issue

V	FU	OP	rd	src 1	src 2	result	exception
0	LSU	LD	x1	x2	x0	128	No
1	MUL	MUL	x2	x2	x2	0xAAAA	No
1	ALU	ADD	x3	x0	x0	16	No
1	ALU	ADD	x2	x3	x1	0xAABA	No

Commit

Issue

Commit

ALU

LSU

MUL

Regfile

Scoreboard – Exception

Decode:

ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)

Issue

V	FU	OP	rd	src 1	src 2	result	exception
1	LSU	LD	x1	x2	x0	128	Yes
0							
1	ALU	ADD	x3	x0	x0	16	No
1	ALU	ADD	x2	x3	x1	0xAABA	No

Commit

Issue

Commit

ALU

LSU

MUL

Regfile

Scoreboard - Commit

Decode:

ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)

Issue



V	FU	OP	rd	src 1	src 2	result	exception
1	LSU	LD	x1	x2	x0	128	Yes
0							
0							
1	ALU	ADD	x2	x3	x1	0xAABA	No

Commit



Issue

Commit

ALU

LSU

MUL

Regfile

Scoreboard - Commit

Decode:

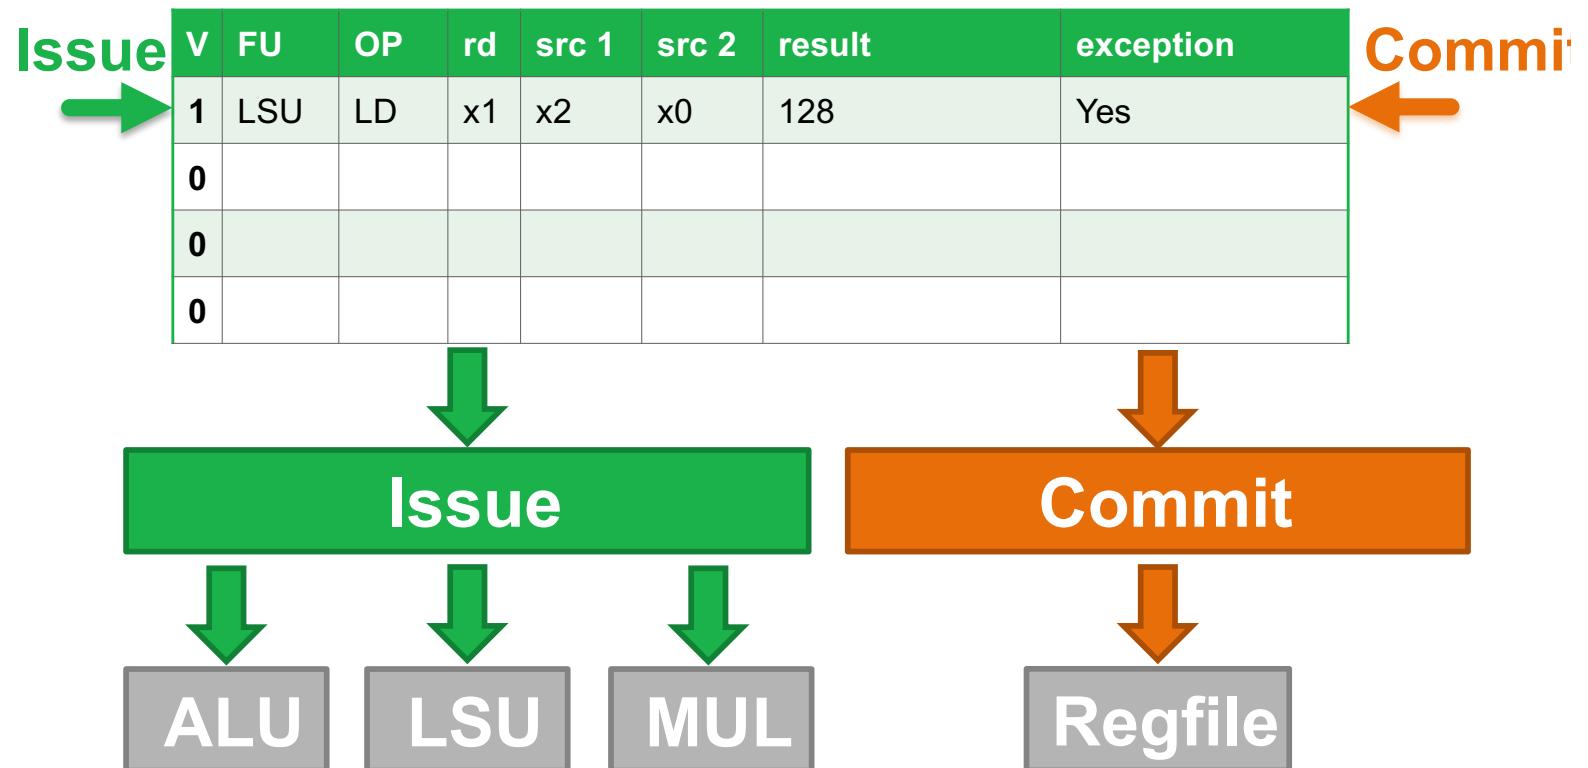
ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)



Scoreboard - Commit

Decode:

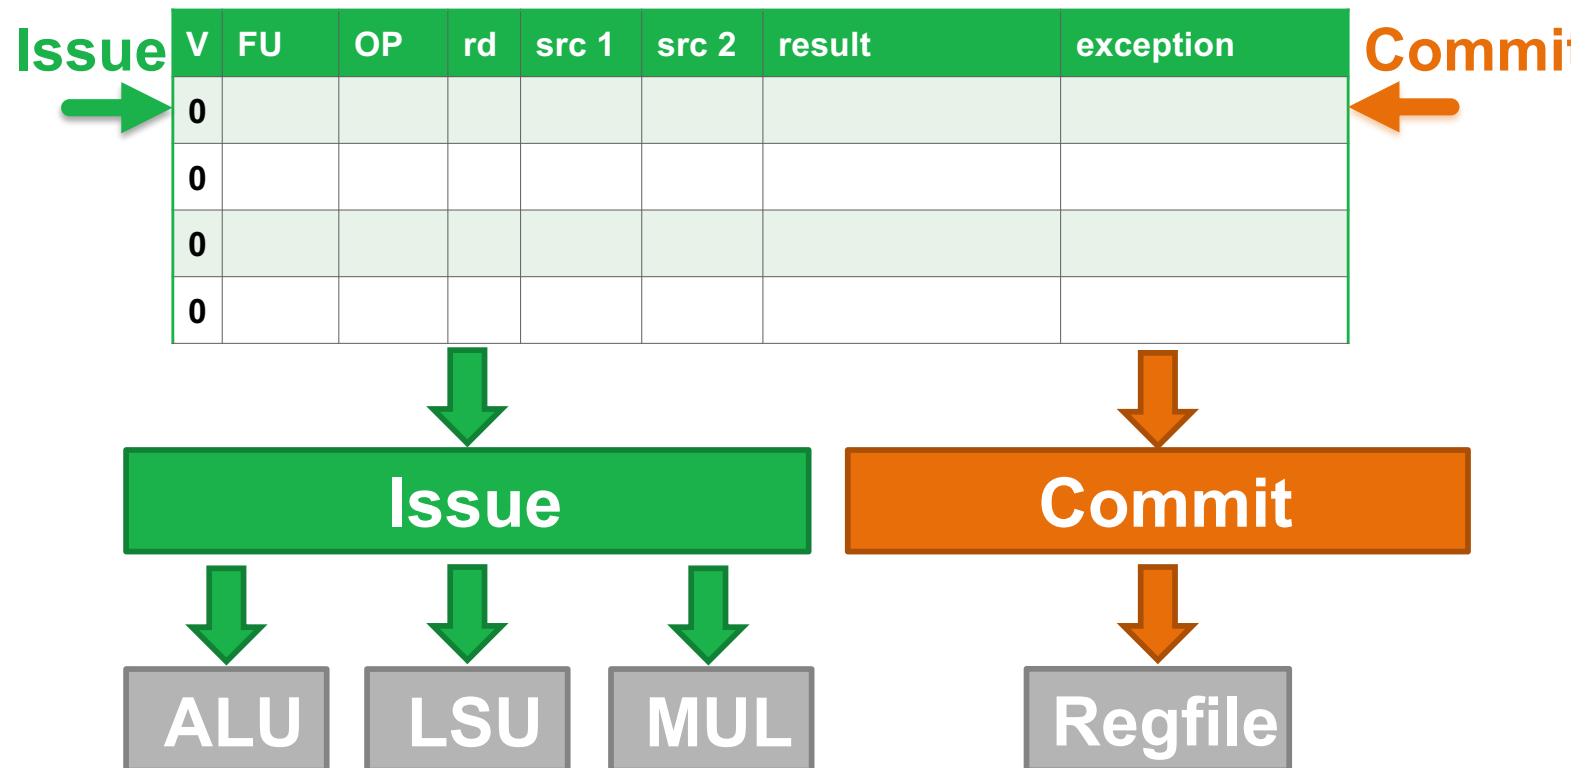
ld x1 \leftarrow x2(0)

mul x2 \leftarrow x2, x2

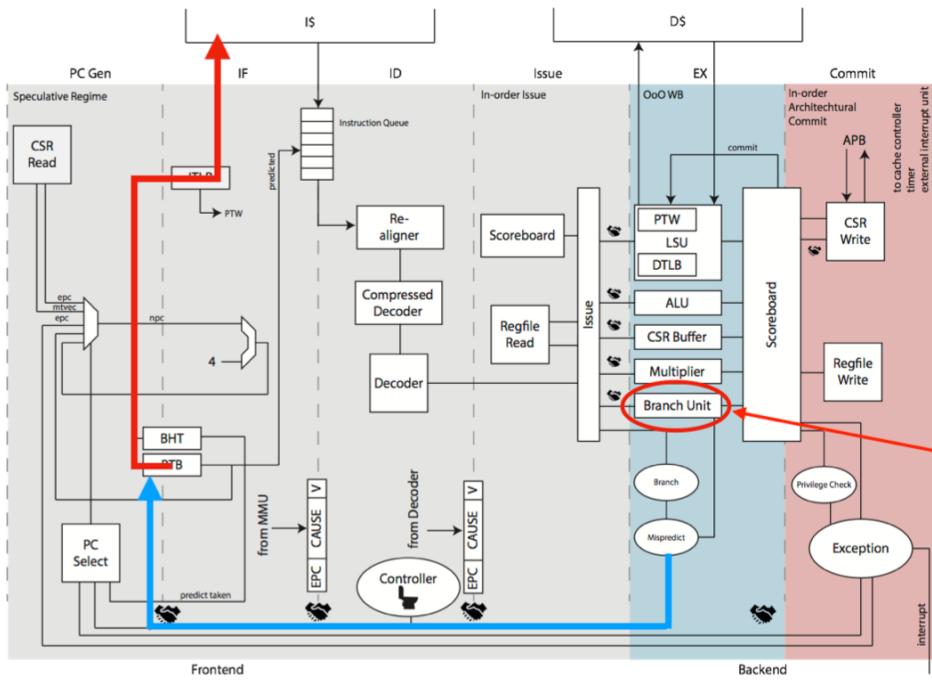
addi x3 \leftarrow x0, 16

addi x2 \leftarrow x3, x1

ld x1 \leftarrow x2(128)



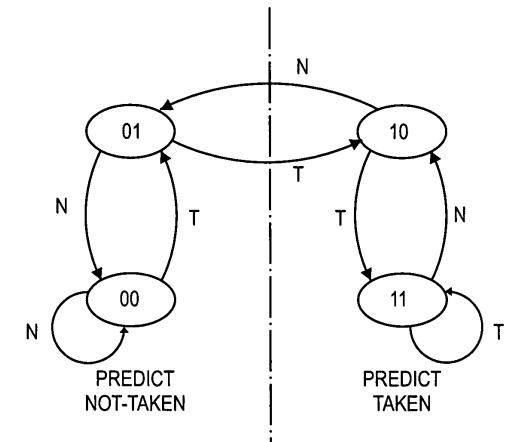
Branch-prediction



Control Hazards

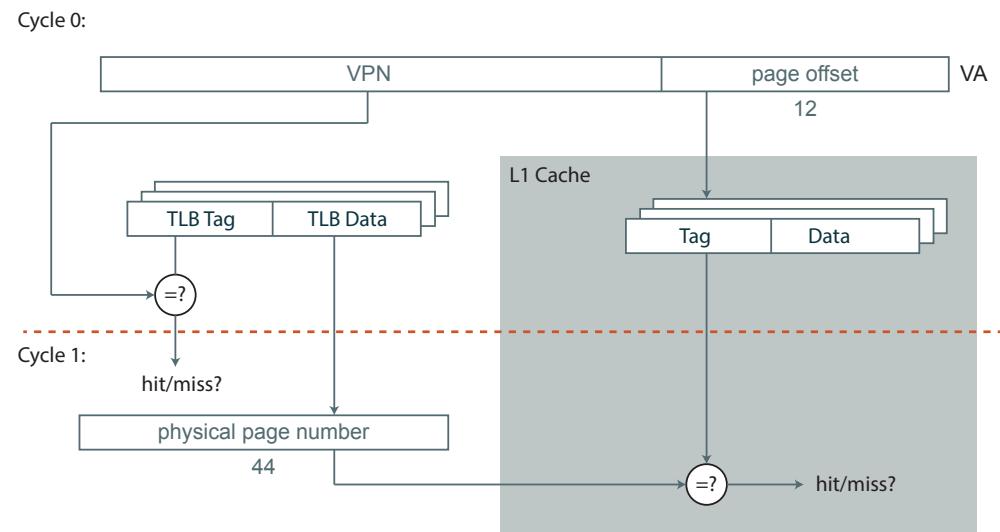
→ Branch Prediction

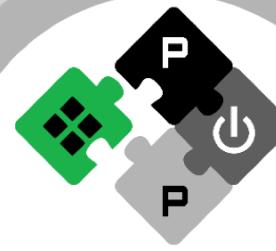
taken?
bneq x4, x6, pc + 16
addi x3, x2, 10



Caches

- Caches are a necessity for larger systems
- Private L1 caches
 - **I\$ (16 kByte, 32 entries, 16 byte cache line, 4 way) - 1,41% MR** (Linux Boot)
 - **D\$ (32 kByte, 32 entries, 16 byte cache line, 8 way) - 3,17% MR** (Linux Boot)
- Currently no (shared) L2 cache
- SRAMs (cache memories) are slow compared to regular logic
- Virtually indexed, physically tagged data cache





PULP
Parallel Ultra Low Power

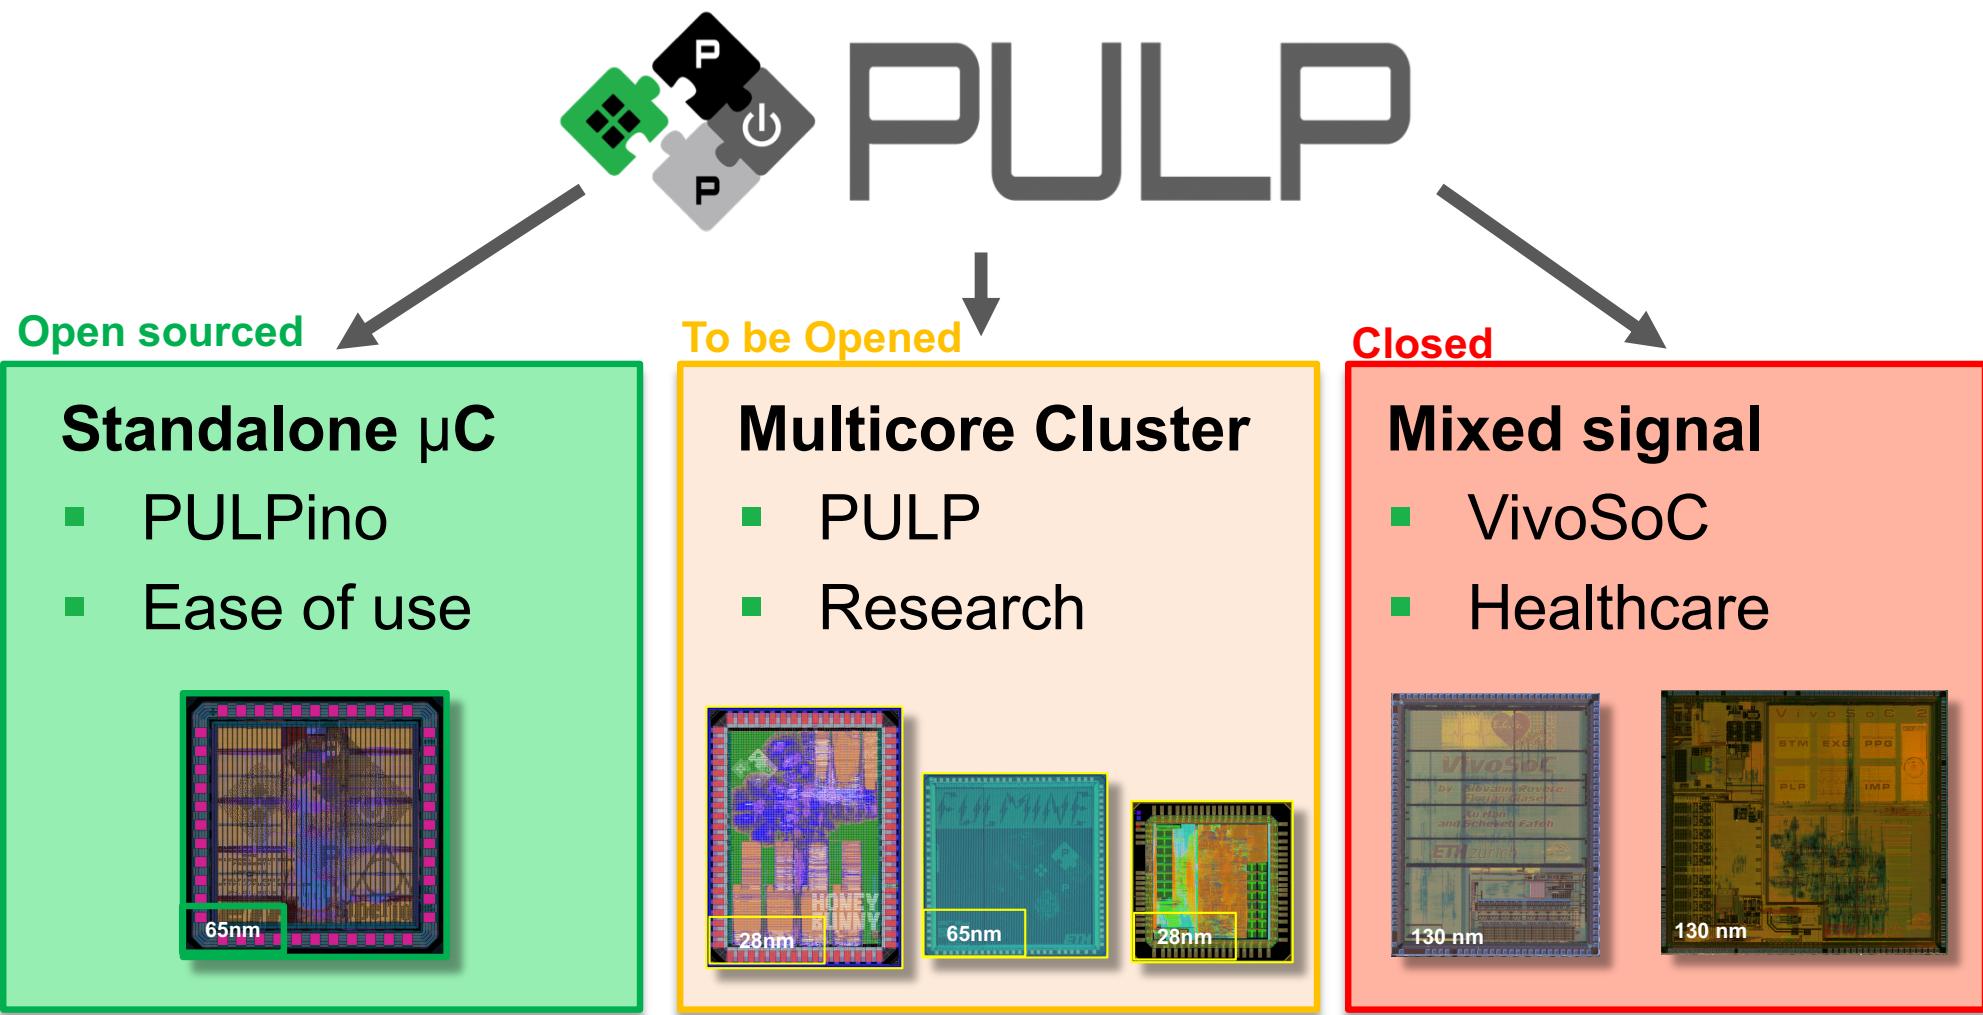


*¹Department of Electrical, Electronic
and Information Engineering*

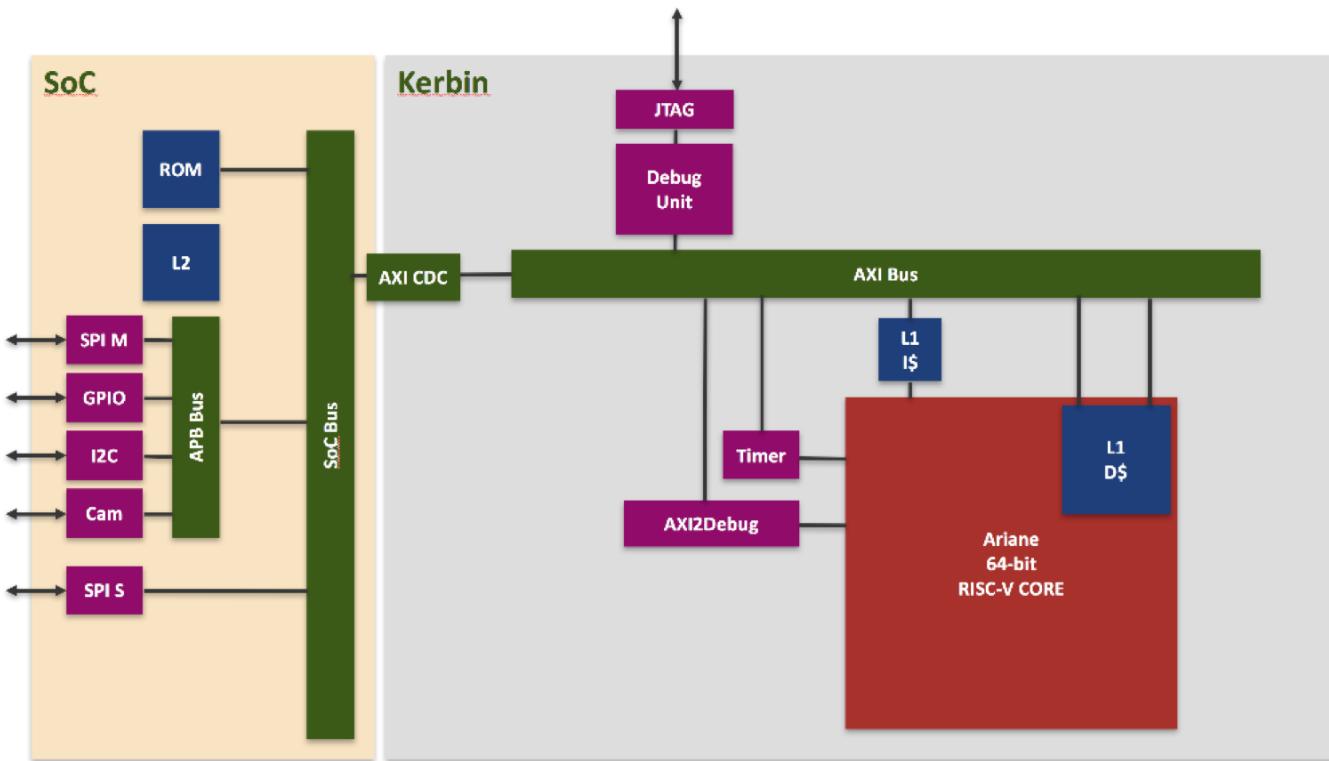
ETH zürich

²Integrated Systems Laboratory

From Cores to (ULP) SoC Platforms



Kerbin: Proof of concept SoC for Ariane



SoC

- PULP Peripherals

CLUSTER

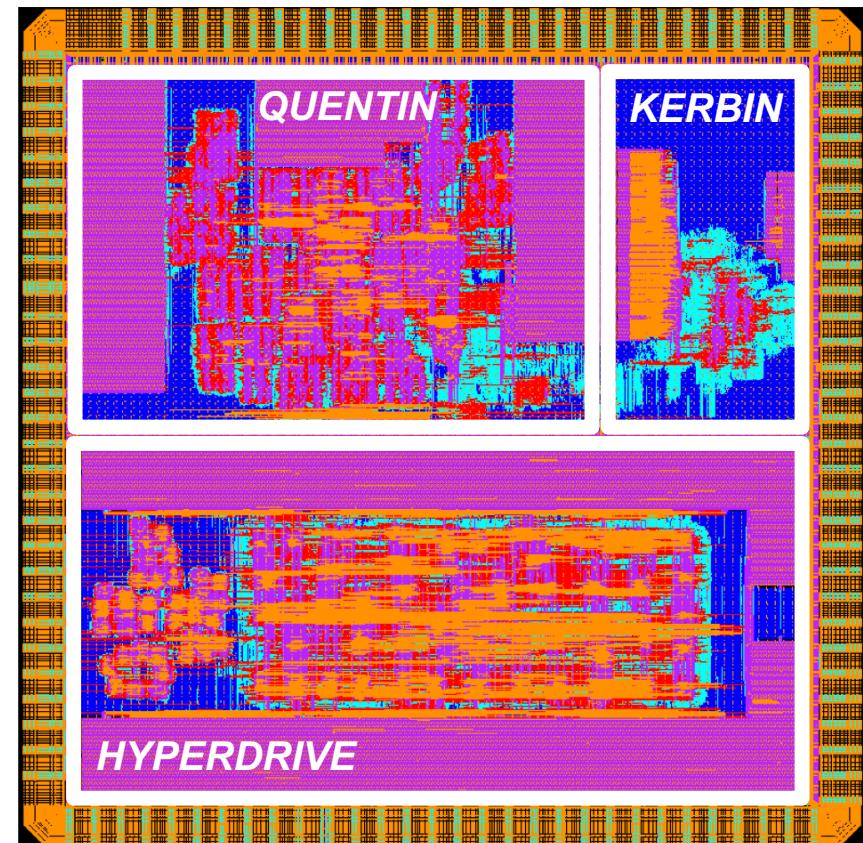
- 64-bit interconnect
- Debug support

→ Taped out in Globalfoundries 22 nm FDX in January

Poseidon

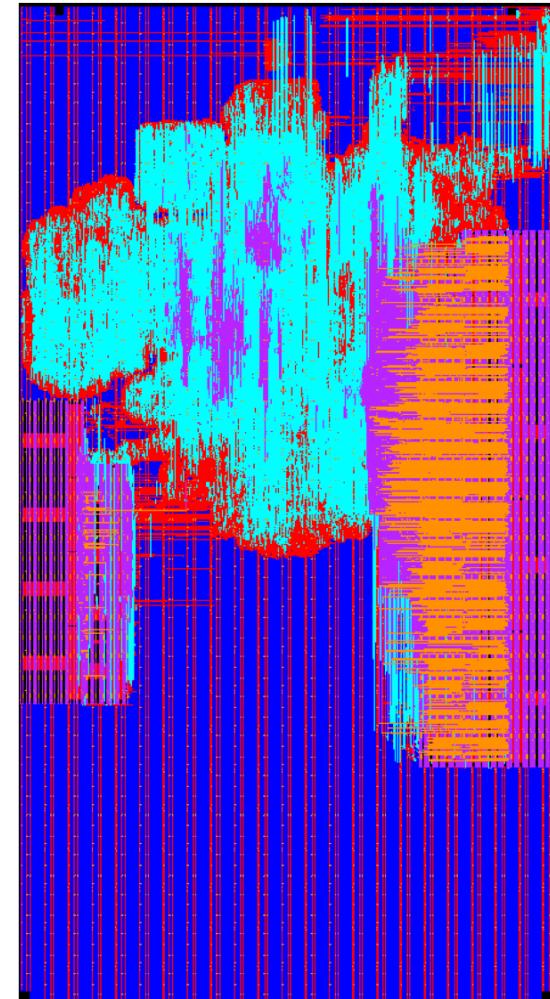
- First 22FDX Testchip of the labs
- Goals:
 - Explore performance/energy efficiency:
 - In the low-power domain → LVT
 - In the high-performance domain → SLVT
 - Explore the voltage range of the technology (LVT + SLVT)
 - Explore body biasing
- Test SoCs:
 - **Poseidon**: ultra-low-power 32-bit microcontroller
 - **Kerbin**: high-performance 64-bit processor
 - **Hyperdrive**: Convolutional Neural Networks Accelerator

Poseidon layout



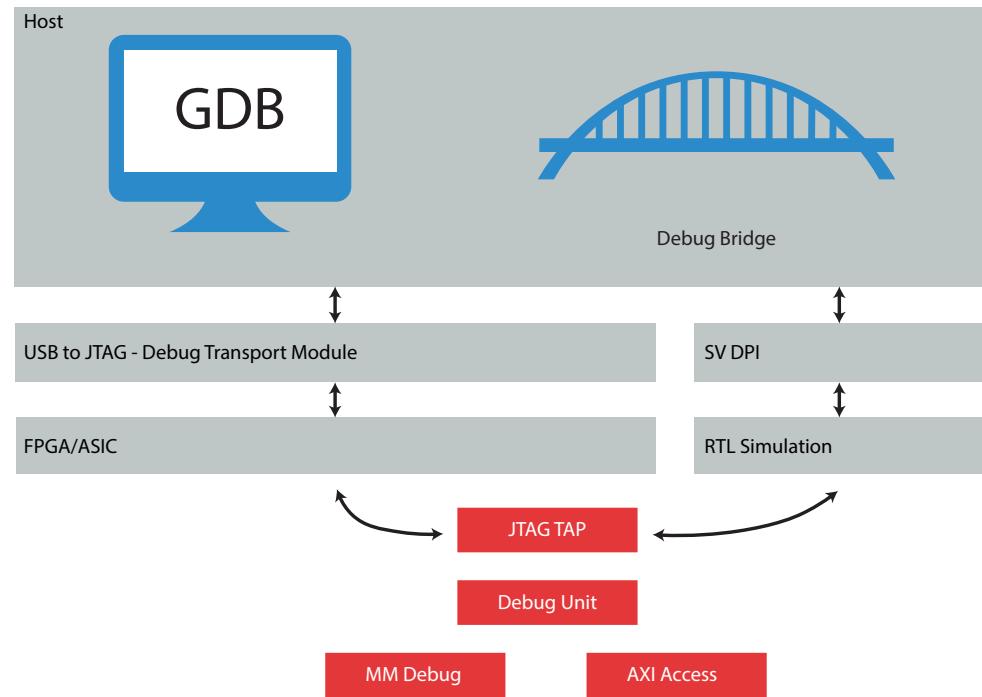
Kerbin Implementation Results

- Silicon implementation in GF22FDX, mixed LVT and SLVT libraries.
- The system features 32 kByte of instruction and 32 kByte of data cache.
 - D\$: 8-way, 4 kByte
 - I\$: 4-way, 8 kByte
- Timing closure: 910 MHz @ SSG, 125/-40 °C, 0.72V – NO BB
- Area: 0.23 mm² – 175 kGE



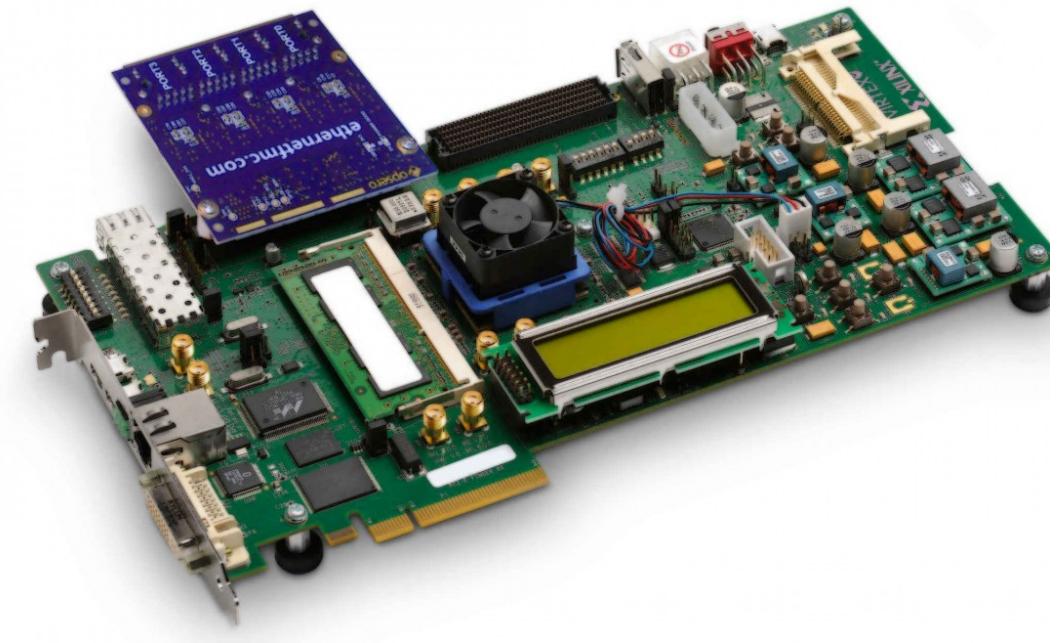
Full Debug support

- Leveraging existing infrastructure:
RISC-V GDB
- Debug Bridge to communicate with hardware
- Allows for:
 - run-control
 - single-step
 - inspection
 - (hardware) breakpoints
- Essential for SW debug and hardware bring-up



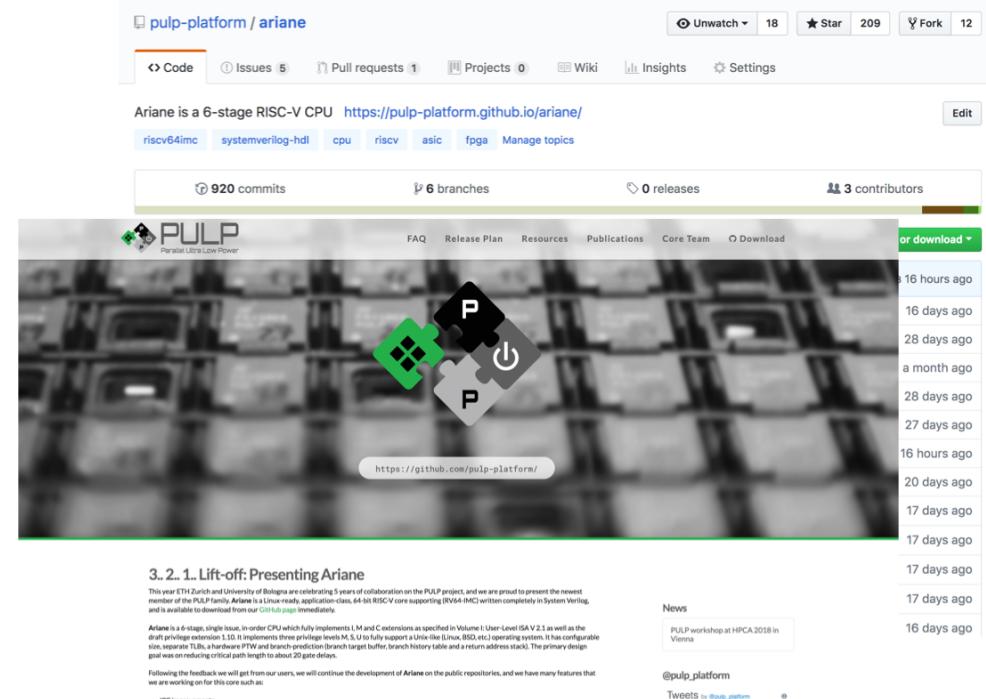
Status Quo

- Full FPGA implementation
 - Xilinx Vertex 7 – VC707
 - **Core:** 50 – 100 MHz
 - **Core:** 15 kLUTs
 - 1 GB DDR3
 - Booting Linux ([Demo](#))
- FPGA implementation allows for fast prototyping – HW/SW codesign
- Area allows for exploration of multi-processor systems



It is Open Source!

- The whole core has been open-sourced
- **OpenSource toolflow for easy exploration:** Verilator
- QuestaSim (commercial simulator) toolflow for more thorough use-cases
- Further development entirely in the open:
 - Further IPC improvements
 - Floating point support
 - Atomic Support
- FPGA mapping following soon!



→ <https://github.com/pulp-platform/ariane>

Structure of this workshop

- Open source hardware and our role (Frank)
- The PULP family tree (Frank)
- Our RISC-V cores: Ariane, RI5CY and friends (Florian)
- Break – Demos
- Accelerators in PULP (Francesco)
- Our Programmable Multi-Core Accelerator – HERO (Andreas)
- Programming PULP (Andreas)

Please interrupt at any time to ask questions

QUESTIONS?

